

Benchmarking Keccak256 proving algorithms: Maru, Axiom and JumpCrypto

September 2023

Abstract

Cryptography is constantly evolving, and with it the consensus algorithms that can be used to keep data private. For the time being SNARK and STARK are the main technologies that become the most popular to research. Each of these algorithms has its own pros&cons, as well as the ways to use them. This paper compares three proof generation implementations for the Keccak256 hash function from Maru, Axiom, and JumpCrypto that implement SNARK/STARK schemes.

We, Maru, offer an implementation of a Keccak256 on STARK, which later "turns" into a SNARK. We also provide the results of a comparison of all three implementations according to such criteria as time for building a circuit, generating & verifying a proof and proof size.

1 Introduction

Zero-knowledge proof is one of the technologies that is currently gaining popularity and does not only transform cryptography, but also improves all existing blockchain infrastructures.

Zero-knowledge proof technology allows proving some information without revealing the information itself. The scheme has two participants in the protocol: the prover and the verifier. The prover claims to have certain information. And the verifier has to check the validity of prover's data. The ZK-SNARK protocol does not require continuous communication or conversation between the prover and verifier. It works based on complex mathematical operations.

ZK-SNARK. The protocol is based on a trusted configuration using such mathematical assumptions as homomorphic functions, blind evaluation, Pinocchio's protocol etc. ZK-SNARK allows users to send transactions on the blockchain in a completely encrypted way. This means that transactions are completely legitimately, but no one can read them. Furthermore, such a transaction cannot be modified by a third party. With ZK-SNARK, we can see that the sender has the funds, but we cannot see how much or where he wants to send it.

ZK-STARK. ZK-STARK provides the ability to share verified data or perform calculations with a third party, without revealing the data to that party. At the same time, it is publicly verifiable. ZK-STARK allows you to verify the banking information of your future business counterparty, without having to disclose your confidential information. The protocol moves computation and storage off the blockchain. Therefore, it improves its scalability and privacy. Services performed outside the blockchain can generate STARK proofs, which simultaneously certify the integrity of off-chain computations. The proofs thus made are then placed back on the blockchain so that any interested party can validate the computation made. This type of zero proof focuses first on scalability and only later on privacy.

SHA3 or Keccak is a hash function consisting of four cryptographic hash functions and two extendable-output functions. These six functions are based on an approach called sponge functions. Sponge functions provide a way to generalize cryptographic hash functions to more general functions with arbitrary-length outputs.

SHA-3 consists of four functional blocks called state function, round constant, buffer function and Keccak function.

The algorithm receives as a matrix, an input called state, which represents an array consisting of bits arranged in a $5 \times 5 \times w$ format, where w is defined as $w = \text{power}(2, l)$. Since we have chosen $l=6$, w is equal to 1600. Therefore, we utilize 1600 bits as the input length S . We denote the bit $(5i + j) \times w + k$ of the input as $a[i][j][k]$. The block permutation function involves twelve plus two times 1 rounds, each comprising of five steps: $\theta(\theta)$, $\rho(\rho)$, $\pi(\pi)$, $\chi(\chi)$, and $\iota(\iota)$.

This paper describes testing of different implementations of Keccak256. The main idea of the work is to demonstrate different approaches and analyze their efficiency.

2 Related implementations

We came across a number of GitHub projects, which implements Keccak256 ZK-proofs:

- Axiom provides SNARK Keccak256 proof using Halo2.
URL: <https://github.com/axiom-crypto/halo2-lib/tree/community-edition/ hashes/zkevm-keccak/src>.
- JumpCrypto offers SNARK Keccak256 proof using plonky2. URL: <https://github.com/JumpCrypto/plonky2-crypto>.
- Maru provides recursive Keccak256 circuit using starky and plonky2 frameworks. The idea is to generate a proof in STARK, that turns to SNARK.
URL: https://github.com/proxima-one/keccak_ctl

Maru implementation of Keccak256 uses Polygon's implementation of Keccak256 for the EVM. In this work, Keccak256 is divided into two entities: sponge

and permutation. To solve the problem of computing the computational integrity of a message, we added a check of the public input, i.e. the hash. After receiving two proofs on STARK, we combine them into one SNARK proof. Such aggregation also allows you to significantly compress the resulting proof.

The benchmark is available on GitHub and anyone can easily contribute by visiting the following link: <https://github.com/proxima-one/keccak-circuit-benchmarks> to understand and test these schemes comprehensively.

Our goal is to show the community the performance and future potential of ZK-technology. As new schemes and frameworks become available, there will be additional opportunities for comparison.

3 Benchmark methodology

In our benchmark we compute the Keccak256 hash function for different data length: $N = 136, 272, \dots, 136\text{KB}$ (except JumpCrypto, because we need more resources there. We take $N = 136, 272, \dots, 19992\text{B}$). We conduct the benchmark of each system using the following performance metrics:

- Circuit building time
- Proof generation time
- Proof verification time
- Proof size in bytes

The chosen parameters are as follows:

- Axiom implementation with input message length ranging from 136 to 100,096 bytes. The parameters used are: `rows_per_round=25`.
- JumpCrypto implementation with input message length ranging from 136 to 19,992 bytes. Large message sizes do not allow accurate measurements.
- Maru implementation with input message length ranging from 136 to 136,000 bytes.

We conducted our benchmark on Xeon E5-2697v4, 128GB DDR4 2400MHz with Linux. The specific version details: Linux Version: Ubuntu 22.04.2 LTS.

4 Benchmark

We will compare the three works by plotting diagrams that include all three implementations.

Building circuit. Given the significant circuit building time in the JumpCrypto implementation, we employed a logarithmic function to scale the data, ensuring that the circuit building graph is more interpretable. The quickest circuit building time was observed with the Axiom implementation, which

recorded a time lower than 28,152 bytes. Following in terms of speed is the JumpCrypto implementation, which outperforms the others for small inputs, up to 816 bytes. The JumpCrypto implementation shows slower performance due to its dependence on the number of blocks for the input message length and curve displayed on the plot was transformed into a logarithmic curve, while the initial dataset exceeds 10744 bytes in just 220 seconds. As a result, the circuit building time is very high. On the other hand, the Maru implementation exhibits better overall performance for larger inputs, surpassing Axiom’s performance from 28,288 bytes onward.

Proof generation. The JumpCrypto implementation yields the best results for proof generation when dealing with small message lengths. However, it’s worth noting that blockchain hashes, particularly those from Ethereum, can reach sizes of up to 10KB. In this context, the Maru implementation demonstrates superior overall performance. The JumpCrypto implementation exhibits faster performance than Axiom for message input sizes up to 2448 bytes. However, as the number of blocks increases, its execution time also increases more rapidly compared to other implementations. On the other hand, the Axiom implementation generally outperforms the JumpCrypto implementation for message sizes starting from 2584 bytes. If your priority is larger message lengths in bytes, then the Maru implementation may be more suitable for your needs. However, it’s crucial to consider your specific requirements carefully. Furthermore, regardless of the message size, the Maru implementation maintains a stable proof generation time that increases gradually compared to the other implementations. In conclusion, the choice of implementation should be tailored to your particular use case and requirements.

Proof size. For a given message length range, the Axiom implementation has the smallest proof size. As for Maru, before aggregation, the proof size for permutations is more than 1,780,000 bytes. But we compress it by turning to SNARK and the final size is ≈ 160 KB. Notably, this compressed size outperforms JumpCrypto for message lengths starting from 2448 bytes. Therefore, if minimizing proof size is a priority, the Axiom implementation would be a favorable choice. However, considering the overall performance and the specific needs of your application, both Maru, JumpCrypto and Axiom offer good results in terms of proof size for the given message length range.

Proof verification. Unlike previous comparisons, the JumpCrypto implementation stands out due to its significantly faster verification time, averaging around 0.0051 seconds. Following closely in terms of verification efficiency is the Axiom implementation, with a lower verification time of approximately 0.148 seconds for small inputs and 0.124 seconds for large inputs, respectively. On the other hand, the Maru implementation exhibits a longer verification time, averaging around 0.187 seconds within the specified range. While there are differences in verification times among the implementations, they may not be noticeable when considering the previous comparisons.

Implementations effectiveness comparing with native verification. The next comparison involves determining the specific point of intersection between the trend lines of each implementation to identify advantage compared to

native check. By plotting a linear trend line for each range, these trend lines can be extended by determining the slope and intercept. When the trend lines intersect, it indicates the effectiveness of zk-implementations compared to native verification. The first "intersecting" trend line belongs to the implementation by JumpCrypto that will cross trend line at around 1,006,462 bytes. The proof verification for this implementation occurs faster than others. Therefore, for overall size of input, the verification will be more efficient for this implementation compared to native verification. Furthermore, the trend lines of native verification and proof verification by Axiom intersect at a size of approximately 34,682,015 bytes. As far as Maru is concerned, the efficiency compared to native checking is the same as the Axiom implementation, and the trendlines will cross at around 41,284,692 bytes.

5 Conclusion

The work demonstrates different approaches and shows their advantages. We conducted the tests of Keccak256 implementations using three different platforms and collected metrics for proof generation and verification time, proof size and specific metrics for each scheme.

JumpCrypto implementation heavily sacrifices circuit building time and proof generation time in order to achieve significant advantage over other competitors in proof verification time. Also for very small messages, JumpCrypto builds circuits and generate proofs faster than some competitors.

Axiom implementation takes more balanced approach, with much faster circuit building, faster proof generation and much smaller proofs but significantly slower verification than JumpCrypto in most cases. In comparison to Maru, this implementation have faster circuit building for smaller messages and slower for larger messages, generally much slower proof generation, much smaller proof sizes and slightly faster verification times.

Maru implementation significantly outperforms JumpCrypto on almost everything, except proof verification time in most cases. In comparison to Axiom, this implementation have comparable circuit building time (Axiom is faster for messages less than about 30 kilobytes, for larger messages Maru is faster), much faster proof generation time, but larger proofs and slightly larger verification time.

By comparing these implementations, we aimed to find the theoretical point at which efficiency favors Zero-knowledge, but this is just a projection. Ultimately, the choice to use these implementations depends on your specific use cases and the criteria that are most important to you.

The crypto community continues to research in search of new frameworks that will bring us closer to efficient data processing. In our opinion, Zero-knowledge proof is a promising direction, the study of which should be continued.

6 Appendix




Figure 1: Comparing building circuit time using **logarithmic** scale function

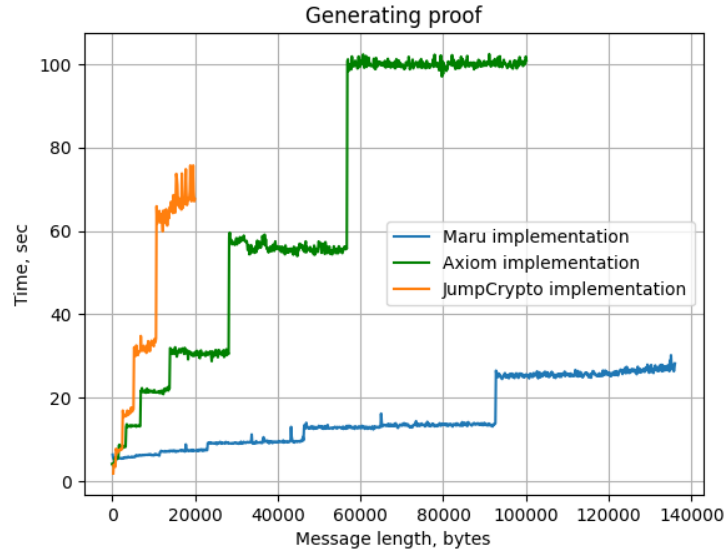


Figure 2: Comparing proof generation time




Figure 3: Comparing proof size




Figure 4: Comparing proof verification time




Figure 5: Comparing effectiveness of native keccak with projected proof verification time of ZK implementations